
ST - Reinforcement Learning: Homework #2
Due:October 8, 2025 @ 23:59

1 Part A: Conceptual Foundation

1.1 How DQN Extends Q-Learning

1.2 The Role of The Target Network

1.3 The Replay Buffer

1.4 Approximation of the Bellman Equation

1.5 MDP versus Deep Q-Learning

2 Part B: DQN Implementation

3 Part C: Double DQN

4 Part D: Double DQN Compared to DQN

5 Part E: Prioritized Experience Replay
